

# Neural Networks and Function Approximation

Krishna Kumar

University of Texas at Austin

*krishnak@utexas.edu*

# Overview

# Outline

# The 1D Poisson Equation: Our Benchmark Problem

Consider the one-dimensional Poisson equation on  $[0, 1]$ :


$$-\frac{d^2 u}{dx^2} = f(x), \quad x \in [0, 1]$$

with boundary conditions:

$$u(0) = 0, \quad u(1) = 0$$


For  $f(x) = \pi^2 \sin(\pi x)$ , the analytical solution is:

$$u(x) = \sin(\pi x)$$



# The Function Approximation Challenge

**Key Question:** Can we learn to approximate  $u(x) = \sin(\pi x)$  from sparse data?




# Traditional vs Neural Network Approaches

## Finite Difference:

- Discretize domain into grid points
- Solve for values at specific locations
- Result: Discrete representation

## Neural Network Approach:


- Learn continuous function  $u_{NN}(x; \theta)$
- Approximate solution over entire domain
- Parameters  $\theta$  trained from sparse data



Finite Difference Approximation

$$\frac{d^2 u}{dx^2} \approx \frac{u_{i+1} - 2u_i + u_{i-1}}{h^2}$$

# Finite Difference vs Neural Networks



Comparison of solution methods and discretization errors

**Key Insight:** Neural networks learn continuous functions, not just discrete values.


# Outline

# The Perceptron: Basic Building Block

A perceptron computes:

$$z = \mathbf{w}^T \mathbf{x} + b = \sum_{i=1}^n w_i x_i + b$$
$$\hat{y} = g(z)$$

where  $g$  is the activation function.



**Key components:** Input vector, weights, bias, activation function, output.

# The Critical Role of Nonlinearity

**Without activation functions:** Multiple linear layers collapse to single linear transformation.

Consider two linear layers:

$$h_1 = W_1x + b_1$$


$$h_2 = W_2h_1 + b_2 = W_2(W_1x + b_1) + b_2$$

$$= (W_2W_1)x + (W_2b_1 + b_2)$$

This is equivalent to:  $h_2 = W_{eq}x + b_{eq}$  (still linear!)

**Problem:** Linear networks can only learn  $y = mx + b$ , but  $\sin(\pi x)$  is curved!

# Demonstrating Linear Network Failure



Linear network cannot approximate  $\sin(\pi x)$

**Conclusion:** Nonlinearity is essential for complex function approximation.

# Common Activation Functions

**Sigmoid:**  $\sigma(x) = \frac{1}{1+e^{-x}}$  (squashes to (0,1))

**Tanh:**  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$  (squashes to (-1,1))

**ReLU:**  $f(x) = \max(0, x)$  (efficient, prevents vanishing gradients)


**Leaky ReLU:**  $f(x) = \max(\alpha x, x)$  (prevents dead neurons)

▶ ReLU Demo

▶ Activation Comparison

# Common Activation Functions

Common Activation Functions (Parameterized)



# Outline

## Can neural networks approximate ANY continuous function?

**Answer: YES! (Cybenko 1989, Hornik 1991)**

A single hidden layer network can approximate any continuous  $f : [0, 1] \rightarrow \mathbb{R}$  to arbitrary accuracy:

$$\|f - F_N\|_\infty < \epsilon \quad \text{where} \quad F_N(x) = \sum_{i=1}^N w_i \sigma(v_i x + b_i) + w_0$$

### Three Big Questions:

- **Why does it work?** Mathematical foundations
- **How does it work?** Constructive proof with ReLU
- **What are the limits?** Width vs depth trade-offs

# Mathematical Foundations: Function Spaces

**Banach Space:** Complete normed vector space - the setting for approximation

## Key Concepts:

- **Vector space:** Add functions, multiply by scalars
- **Norm:** Measure "distance"
- **Complete:** Limits exist
- **Dense:** Can get arbitrarily close

## Common Norms:


- $\|f\|_\infty = \max|f(x)|$
- $\|f\|_2 = \sqrt{\int f^2 dx}$
- $\|f\|_1 = \int |f| dx$

**Hilbert Space:** Banach space with inner product  $\langle f, g \rangle = \int f \cdot g dx$

- Enables orthogonality and projections
- Basis expansions: Fourier series, wavelets

**UAT Statement:** Neural networks are **dense** in  $C([0, 1])$  under  $\|\cdot\|_\infty$

# Density Illustrated: Approximating $\sin(\pi x)$



figs/uat-approximation-progression.png

# Constructive Proof: ReLU Decomposition

**How to build ANY function from ReLU units:**

figs/relu-decomposition-sinpi.png

# Interactive ReLU Builder


figs/relu-interactive-demo.png

# Bias Terms: The Breakpoint Controllers

**Critical Insight:** Bias determines WHERE ReLU "breaks"

For ReLU neuron:  $h = \max(0, wx + b)$

- Activates when:  $wx + b = 0$
- **Breakpoint:**  $x = -b/w$



# Proof by Contradiction: The Beautiful Argument

**Assume:** NNs cannot approximate  $\sin(\pi x)$  within  $\epsilon$

**The Logic Chain:**

1. Gap exists  $\Rightarrow$  detector exists
2. Detector = linear functional  $L$
3.  $L = \text{measure } \mu$  (Riesz)
4.  $\mu$  annihilates all sigmoids
5. Sigmoids  $\rightarrow$  half-spaces
6. Half-spaces isolate points
7.  $\mu = 0$  everywhere
8. But  $L(\sin(\pi x)) \neq 0$
9. **Contradiction!**




figs/contradiction-visual

The impossible detector

**Conclusion:** Our assumption is false. NNs CAN approximate ANY continuous function!

# The Half-Space Argument Visualized



figs/halfspace-isolation.png

# Activation Function Comparison


**Not all activations are universal approximators!**

figs/activation-comparison.png

# Why Parabolic Fails: A Special Case

figs/parabolic-failure.png

# Experimental Verification: Width vs Error



figs/width-vs-error.png

# Sobolev Spaces: Approximating Derivatives

**UAT extends to derivatives!**

figs/derivative-approximation.png

# Width vs Depth: The Trade-off

**UAT guarantees ONE layer works, but depth is more efficient!**

figs/shallow-vs-deep-sin100x.png

# The Limits of UAT

## What UAT tells us:

- ✓ Approximation is POSSIBLE
- ✓ One hidden layer is sufficient
- ✓ Works for ANY continuous function

## What UAT doesn't tell us:

- How many neurons needed (could be millions!)
- How to find the weights (optimization)
- Generalization beyond training data
- Computational efficiency

## Practical Implications:

- Deep networks often more efficient
- Problem-specific architectures help
- Regularization needed for generalization
- UAT is a "permission slip" - not a recipe!

## Why UAT matters for SciML:

### ① PDE Solutions are Continuous

- Heat equation, wave equation, Navier-Stokes
- UAT guarantees NNs can represent solutions

### ② Mesh-Free Approximation

- No grid needed
- Continuous representation
- Evaluate anywhere in domain

### ③ Automatic Differentiation

- Exact derivatives for free
- Critical for PDE residuals
- Sobolev space approximation

### ④ High Dimensions

- UAT holds for  $\mathbb{R}^n \rightarrow \mathbb{R}^m$
- Avoids curse of dimensionality (sometimes)

**The Bottom Line:** UAT provides theoretical foundation for PINNs

# Summary: Universal Approximation Theorem

## Mathematical Foundations:

- Banach/Hilbert spaces
- Dense subsets
- Function norms

## Two Proofs:

- Constructive (ReLU)
- Contradiction (Hahn-Banach)

## Key Insights:

- Bias = breakpoint control
- Not all activations universal
- Width vs depth trade-offs

## Interactive Resources:

- UAT Demo
- Theory Notebook
- Perceptron Basics

## Applications:

- Physics-informed NNs
- Function approximation
- PDE solving
- Scientific computing

**UAT: The theoretical permission slip for neural networks in science!**

# Single Hidden Layer Architecture


For 1D input  $x$ , a single-layer network with  $N_h$  hidden neurons:

$$\mathbf{z}^{(1)} = W^{(1)}x + \mathbf{b}^{(1)} \quad (\text{pre-activation})$$

$$\mathbf{h} = g(\mathbf{z}^{(1)}) \quad (\text{hidden layer output})$$

$$z^{(2)} = W^{(2)}\mathbf{h} + b^{(2)} \quad (\text{output layer})$$

$$\hat{y} = z^{(2)} \quad (\text{final prediction})$$



Single hidden layer neural network

# Outline

# Training Process

**Goal:** Find optimal parameters  $\theta^*$  that minimize loss function.  
Loss function (MSE):

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (u_{NN}(x_i; \theta) - u_i)^2$$

Optimization problem:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta)$$

## Training Steps:

- ① Forward pass: compute predictions
- ② Calculate loss
- ③ Backward pass: compute gradients
- ④ Update parameters
- ⑤ Repeat until convergence


# The Gradient Problem

Training neural networks requires gradients:  $\frac{\partial \mathcal{L}}{\partial \theta}$  for all parameters  $\theta$ .

Traditional approaches have fundamental flaws:

- **Manual:** Exact, but error-prone and doesn't scale
- **Symbolic:** Exact, but "expression swell"
- **Numerical:**  
 $f'(x) \approx \frac{f(x+h) - f(x-h)}{2h}$ 
  - Inaccurate (rounding errors)
  - Expensive (multiple evaluations)

## DIFFERENTIATION



For a neural network with thousands of parameters:

$$\theta = \{W^{(1)}, \mathbf{b}^{(1)}, W^{(2)}, \mathbf{b}^{(2)}, \dots\}$$

We need a fourth approach: Automatic Differentiation!

# Automatic Differentiation: The Solution

Every function is a computational graph of elementary operations.

Consider:  $y = x_1^2 + x_2$

**Evaluation Trace:**

- ①  $v_1 = x_1^2$
- ②  $y = v_1 + x_2$


$$x_1 = 2, x_2 = 3$$

$$y = 7$$

$$y = x_1^2 + x_2$$

$$y = 2^2 + 3$$


$$y = 7$$



This decomposition is the key that makes AD possible. By applying the chain rule to each elementary step, we can compute exact derivatives.

# Forward Mode Automatic Differentiation

Forward mode AD propagates derivative information **forward** through the graph, alongside the function evaluation.



▶ AD Demo

# Forward Mode: Step-by-Step Example

Let's compute  $\frac{\partial y}{\partial x_1}$  for  $y = x_1^2 + x_2$ .

## 1. Seed the input w.r.t. $x_1$

Set  $\dot{x}_1 = 1$  and  $\dot{x}_2 = 0$ .

## 2. Forward Propagation

- $v_1 = x_1^2 \implies \dot{v}_1 = 2x_1 \cdot \dot{x}_1 = 2x_1 \cdot 1 = 2x_1$
- $y = v_1 + x_2 \implies \dot{y} = \dot{v}_1 + \dot{x}_2 = 2x_1 + 0 = 2x_1$

## Result

The final propagated tangent  $\dot{y}$  is the derivative:  $\frac{\partial y}{\partial x_1} = 2x_1$ .

**Key Idea:** To get the derivative w.r.t.  $x_2$ , we would need a *new pass* with seeds  $\dot{x}_1 = 0, \dot{x}_2 = 1$ .

# Reverse Mode Automatic Differentiation


Reverse mode AD (or **backpropagation**) computes derivatives by propagating information **backward** from the output.

## Algorithm:

- ① **Forward Pass:** Evaluate the function and store all intermediate values.
- ② **Backward Pass:** Starting from the output, use the chain rule to propagate "adjoints" ( $\bar{v} = \frac{\partial y}{\partial v}$ ) backward through the graph.

# Reverse Mode: One Pass for All Gradients

The backward pass systematically applies the chain rule.  
Let's trace the computation for  $y = x_1^2 + x_2$ :



$$y = x_1^2 + x_2$$

$$\nabla y = \left( \frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2} \right)$$

$$\nabla y = (2x_1, 1)$$

We can find the gradient with respect to the output  $y$

Using  $x_1 = 2, x_2 = 3$ :

$$\nabla y = (4, 1)$$

Reverse mode computes **all** partial derivatives in a single backward pass.

# When to Use Forward vs. Reverse Mode

The choice depends on the dimensions of your function  $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ .

## Forward Mode

- **Cost:** One pass per *input* variable.
- **Efficient when:** Few inputs, many outputs ( $n \ll m$ ).
- **Computes:** Jacobian-vector products ( $J \cdot v$ ).

## Reverse Mode

- **Cost:** One pass per *output* variable.
- **Efficient when:** Many inputs, few outputs ( $n \gg m$ ).
- **Computes:** Vector-Jacobian products ( $v^T \cdot J$ ).

**Machine Learning context:** We have millions of parameters (inputs,  $n$ ) and a single scalar loss function (output,  $m = 1$ ).

**Reverse mode is the clear winner!**


# Gradient Descent Algorithm

## Basic Algorithm:

- ① Initialize weights randomly
- ② Loop until convergence:
  - ③ Compute gradient  $\frac{\partial \mathcal{L}}{\partial \theta}$
  - ④ Update weights:  
$$\theta \leftarrow \theta - \eta \frac{\partial \mathcal{L}}{\partial \theta}$$
  - ⑤ Return weights

$$\theta_{new} = \theta_{old} - \eta \nabla \mathcal{L}(\theta)$$

where  $\eta$  is the learning rate.




Gradient descent optimization

▶ SGD Demo

# Width vs Approximation Quality


**Hypothesis:** More neurons  $\rightarrow$  better approximation

**Experiment:** Train networks with 5, 10, 20, 50 neurons




Approximation quality vs network width

# Training/Validation



# Training Convergence Analysis




Training convergence and final loss vs width

## Key Findings:

- 50 neurons:  $\sim 10\times$  better than 5 neurons
- Logarithmic improvement with width
- Convergence rate similar across widths

# Overfitting and Model Capacity

**Problem:** High-capacity networks can memorize training data.




Under and overfitting illustration

**Detection:** Monitor validation loss during training.

**Solutions:** More data, regularization, early stopping, simpler architectures.

# Demonstrating Overfitting



Normal (50 neurons) vs very wide (500 neurons) network

**Observation:** Very wide networks may generalize worse despite lower training loss.


# Outline

# Bias Terms: The Hidden Controllers

**Key Insight:** In ReLU networks, bias terms determine breakpoints!

For neuron  $i$ :  $h_i = \max(0, w_i x + b_i)$

- Activates when:  $w_i x + b_i = 0$
- Breakpoint at:  $x = -b_i / w_i$
- Bias controls position, weight controls slope





# The XOR Problem: Historical Crisis

## XOR Truth Table:

$x_1$	$x_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	0

**The Crisis:** No single line can separate these classes!



XOR is not linearly separable

# True Single-Layer vs Multi-Layer

## Critical Distinction:

- **True Single-Layer:** Input → Output (NO hidden layers)
- **Multi-Layer:** Input → Hidden → Output (1+ hidden layers)

## True Single-Layer (fails):

$$y = \sigma(w_1x_1 + w_2x_2 + b)$$


## Multi-Layer (succeeds):

$$h_1 = \sigma(w_{11}x_1 + w_{12}x_2 + b_1)$$

$$h_2 = \sigma(w_{21}x_1 + w_{22}x_2 + b_2)$$

$$y = \sigma(v_1h_1 + v_2h_2 + b_3)$$

# XOR Solution: Decision Boundaries



Linear (fails) vs Non-linear (succeeds) decision boundaries

**Key Insight:** Hidden layers enable curved boundaries through non-linear transformations.

## XOR Decomposition:

$$h_1 = \sigma(\text{weights}) \quad (\approx \text{OR gate})$$

$$h_2 = \sigma(\text{weights}) \quad (\approx \text{AND gate})$$


$$y = \sigma(v_1 h_1 + v_2 h_2 + b_3) \quad (\approx \text{OR AND NOT})$$

**Result:** XOR = (OR) AND (NOT AND) = compositional solution!

**General Principle:** Complex functions = composition of simple functions.

# Beyond XOR: High-Frequency Functions

**Test Case:**  $f(x) = \sin(\pi x) + 0.3 \sin(10\pi x)$



Shallow (100 neurons) vs Deep (4 layers, 25 each) networks

**Result:** Deep networks achieve better performance with fewer parameters.

# Historical Timeline: From Crisis to Revolution

Year	Event	Impact
1943	McCulloch-Pitts neuron	Foundation laid
1957	Rosenblatt's Perceptron	First learning success
<b>1969</b>	<b>Minsky &amp; Papert: XOR</b>	<b>Showed limits</b>
1970s-80s	"AI Winter"	Funding dried up
1986	Backpropagation	Enabled multi-layer training
1989	Universal Approximation	Theoretical foundation
2006+	Deep Learning Revolution	Depth proves essential


**Lesson:** XOR taught us that depth is necessity, not luxury.

# UAT: Width vs Depth Trade-offs

## Universal Approximation guarantees:

- ONE hidden layer CAN approximate any continuous function
- But may need exponentially many neurons

## Depth advantage for high-frequency functions:



# Four Key Insights on Depth

## 1. Representation Efficiency

- Shallow: May need exponentially many neurons
- Deep: Hierarchical composition is exponentially more efficient

## 2. Feature Hierarchy

- Layer 1: Simple features (edges, patterns)
- Layer 2: Feature combinations (corners, textures)
- Layer 3+: Complex abstractions (objects, concepts)

## 3. Geometric Transformation

- Each layer performs coordinate transformation
- Deep networks "unfold" complex data manifolds

## 4. Compositional Learning

- Complex functions = composition of simple functions
- Build complexity incrementally

# Summary and Conclusions

## Key Takeaways:

- Neural networks learn continuous functions from sparse data
- Nonlinearity is essential for complex function approximation
- Universal Approximation Theorem provides theoretical foundation
- UAT proven via construction (ReLU) and contradiction (Hahn-Banach)
- Bias terms in ReLU = breakpoints in piecewise approximation
- Width increases capacity but depth is more efficient
- Historical XOR problem revealed importance of hidden layers
- Deep networks enable hierarchical feature learning

## Applications:

- Scientific computing and PDE solving (solutions are continuous functions)
- Function approximation and regression
- Pattern recognition and classification
- Physics-informed neural networks (PINNs)
- Sobolev spaces: Can approximate derivatives too!

Questions?

Thank you!

**Contact:**  
Krishna Kumar  
*krishnak@utexas.edu*  
University of Texas at Austin