

# Neural Ordinary Differential Equations

Krishna Kumar

University of Texas at Austin

*krishnak@utexas.edu*

# Overview

1 From ResNets to Continuous Dynamics

2 NeuralODE Architecture

# Learning Objectives

- Understand the connection between ResNets and continuous dynamics
- Master the Neural ODE framework and adjoint method
- Implement ODENets for image classification
- Apply continuous normalizing flows for generative modeling
- Build latent ODE models for irregular time series

▶ Open Notebook

# The ResNet Formula

A residual network transforms hidden states layer by layer:

$$h_{t+1} = h_t + f(h_t, \theta_t) \quad (1)$$

where  $t \in \{0, 1, \dots, T\}$  indexes the layers.

## The Key Question

What happens as we add more layers ( $T \rightarrow \infty$ ) and take smaller steps?

# The Euler Connection

The ResNet update is the **Euler discretization** of an ODE:

$$\frac{dh(t)}{dt} = f(h(t), t, \theta) \quad (2)$$

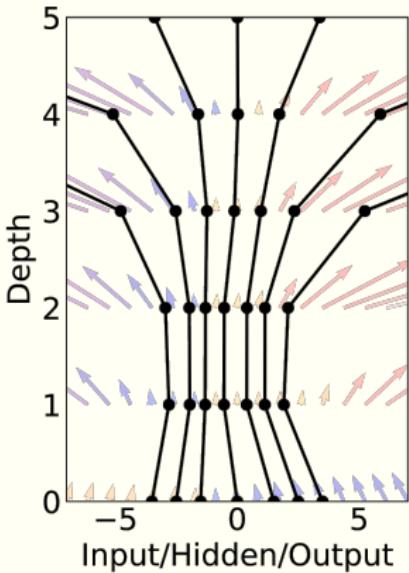
## Key Insight

A ResNet with infinitely many infinitesimal layers  $\equiv$  solving an ODE

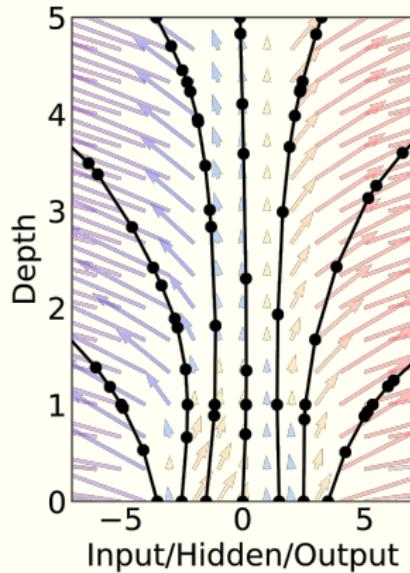
Instead of specifying discrete layers, we parameterize the **derivative** of the hidden state using a neural network.

# ResNet vs ODE

Residual Network



ODE Network



## Residual Network

- Discrete transformations
- Fixed number of layers  $T$
- $h_{t+1} = h_t + f(h_t)$
- Memory:  $\mathcal{O}(T)$

## Neural ODE

- Continuous dynamics
- Adaptive depth
- $\frac{dh}{dt} = f(h(t), t)$
- Memory:  $\mathcal{O}(1)$

# Basic Neural ODE Components

## 1. ODE Function $f(h, t, \theta)$

Neural network that computes the derivative  $\frac{dh}{dt}$

## 2. ODE Solver

Integrates  $h(t)$  from  $t_0$  to  $t_1$ :

$$h(t_1) = h(t_0) + \int_{t_0}^{t_1} f(h(t), t, \theta) dt \quad (3)$$

## 3. Adjoint Method

Computes gradients  $\frac{\partial L}{\partial \theta}$  efficiently

# Memory Efficiency: The Adjoint Method

## Standard Backpropagation:

- Store all intermediate layer activations
- Memory:  $\mathcal{O}(L)$  where  $L = \text{number of layers}$

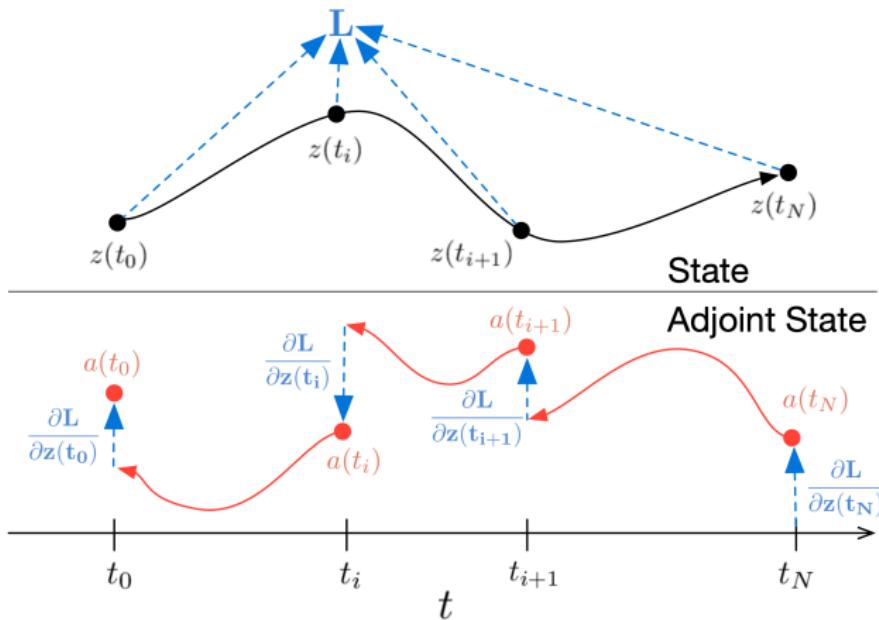
## Adjoint Method:

- Solve a second ODE backwards in time
- Recompute forward states during backward pass
- Memory:  $\mathcal{O}(1)$  independent of depth

The adjoint state  $a(t) = \frac{\partial L}{\partial h(t)}$  evolves as:

$$\frac{da(t)}{dt} = -a(t)^T \frac{\partial f(h(t), t, \theta)}{\partial h} \quad (4)$$

# The Adjoint Method Visualized



- **Forward:** Solve ODE from  $t_0$  to  $t_1$
- **Backward:** Solve augmented ODE from  $t_1$  to  $t_0$
- Automatically handled by `odeint_adjoint`

# Hyperparameter Selection

## ODE Solver Tolerance

- `rtol, atol`: Control accuracy
- Higher tolerance → faster but less accurate
- Typical: `rtol=1e-3, atol=1e-4`

## Solver Method

- **Adaptive**: '`dopri5`', '`adams`' (recommended)
- **Fixed-step**: '`euler`', '`rk4`' (for debugging)

## Integration Time

- Usually  $T = 1.0$  (can be learned)
- Longer  $T \rightarrow$  more expressive but slower

# Extensions

## Augmented Neural ODEs

Add extra dimensions to avoid topological constraints

## Second-order Neural ODEs

Include acceleration:  $\frac{d^2 h}{dt^2} = f(h, \frac{dh}{dt}, t)$

## Stochastic Differential Equations (SDEs)

Add noise for uncertainty:  $dh = f(h, t)dt + g(h, t)dW$

## Hamiltonian Neural Networks

Preserve energy and symplectic structure

# Summary

## Key Takeaways

- ① Neural ODEs = continuous-depth neural networks
- ② ResNets → ODEs via Euler discretization
- ③ Adjoint method enables  $\mathcal{O}(1)$  memory training
- ④ Applications: classification, generative models, time series

## The Big Idea

Parameterize the **derivative** of hidden states, not the states themselves

# References

-  Chen, R. T. Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2018).  
Neural Ordinary Differential Equations.  
*NeurIPS 2018* (Best Paper Award).
-  Grathwohl, W., Chen, R. T. Q., Bettencourt, J., Sutskever, I., & Duvenaud, D. (2019).  
FFJORD: Free-form Continuous Dynamics for Scalable Reversible Generative Models.  
*ICLR 2019*.
-  Rubanova, Y., Chen, R. T. Q., & Duvenaud, D. (2019).  
Latent ODEs for Irregularly-Sampled Time Series.  
*NeurIPS 2019*.